

# Switching Linear Dynamical Systems

Information Theory and Inference

# Intuition

## Switching state space models

The idea behind switching state space models is quite intuitive and can be summed up as

“The macroscopic behavior of a system depends on hidden states...”

One can be seen when driving a vehicle. The macroscopic motion is essentially governed by four basic behaviors (accelerate, brake, turn left, turn right)

# Base Model

## The layers


The model is composed of:

$K$  discrete latent states

$M$  continuous latent states

$N$  visible states

Finally, the time step is denoted with  $t = 1, \dots, T$



# Base Model

## The dynamics

The latent discrete dynamic is Markovian


$$| \mathbf{z}_t \sim \pi_{\mathbf{z}_{t-1}} \in \mathbb{R}^{K \times K}$$

And it is propagated through the other layers as

$$| \mathbf{x}_t = \mathbf{A}_{\mathbf{z}_t} \mathbf{x}_{t-1} + \mathbf{b}_{\mathbf{z}_t} + \mathbf{v}_t$$

$$| \mathbf{y}_t = \mathbf{C}_{\mathbf{z}_t} \mathbf{x}_t + \mathbf{d}_{\mathbf{z}_t} + \mathbf{w}_t$$

Where  $\mathbf{v}_t, \mathbf{w}_t$  are Gaussian noises



# Base Model

## Hidden continuous layer

Remember how

$$x_t = A_{z_t} x_{t-1} + b_{z_t} + v_t \quad \text{with} \quad v_t \sim \mathcal{N}(0, Q_{z_{t+1}})$$

It follows that

$$A_k \in \mathbb{R}^{M \times M}$$

$$b_k \in \mathbb{R}^M$$

$$Q_k \in \mathbb{R}^{M \times M}$$

In principle different for any latent state  $k = 1, \dots, K$

# Base Model

## Visible layer

In a similar way

$$y_t = C_{z_t} x_t + d_{z_t} + w_t \quad \text{with} \quad w_t \sim \mathcal{N}(0, S_{z_{t+1}})$$

It follows that

$$C_k \in \mathbb{R}^{N \times M}$$


$$d_k \in \mathbb{R}^N$$

$$Q_k \in \mathbb{R}^{N \times N}$$

Also different for any latent state  $k = 1, \dots, K$

# Base Model

## Recap



$$\mathcal{D} = [(\pi_k, A_k, Q_k, b_k, C_k, S_k, d_k)]_{k=1}^K$$

# Simplified Model

## SMH my head



Despite having developed the model from a theoretical standpoint with three layers in practice we only used two of them, the two main reasons were

- 1 - Explainability
- 2 - Link to problem complexity



# Final Model

## Simplified model




# Final Model

## Simplified model

$$z_t \sim \pi_{z_{t-1}}$$

$$y_t = A_{z_t} y_{t-1} + b_{z_t} + v_t$$



# Numerical Implementation

**JAGS vs Stan**

**MARTYN  
PLUMMER**



**VS**

**BOB  
CARPENTER**



# Numerical Implementation

## Stan - The inference problem

The first approach was setting up an appropriate Hamiltonian Monte Carlo model with Rstan.

Once the model was set, we needed a way to actually retrieve the latent states  $z$ . This was performed in two steps

- 1 - Likelihood estimation with the *Forward Algorithm*
- 2 - MAP  $z$  sequence retrieval with the *Viterbi Algorithm*

# Numerical Implementation

## Forward algorithm

The forward algorithm has been implemented for both sampling efficiency and a Stan requirement, since discrete parameters are not allowed

The final goal is to compute the joint likelihood

$$p(x_{1:T}, y_{1:T})$$

We can start from

$$\gamma_t(k) = p(z_t = k, x_{1:t}, y_{1:t})$$

# Numerical Implementation

## Forward algorithm

From

$$\gamma_t(k) = p(z_t = k, x_{1:t}, y_{1:t})$$

We can sum over the hidden states  $z$  at  $t - 1$  and then apply the chain rule to retrieve a recurrent expression

$$\gamma_t(k) = p(y_t | z_t = k, x_t) p(x_t | z_t = k, x_{t-1}) \sum_{j=1}^K \pi_{jk} \gamma_{t-1}(j)$$

# Numerical Implementation

## Forward algorithm

Finally

$$\mathcal{L}_k(y_t) = p(y_t | z_t = k, x_t) = \mathcal{N}(C_k x_t + d_k, S_k)$$

$$\mathcal{L}_k(x_t) = p(x_t | z_t = k, x_{t-1}) = \mathcal{N}(A_k x_{t-1}, b_k, Q_k)$$

And so

$$\gamma_t(k) = \mathcal{L}_k(y_t) \mathcal{L}_k(x_t) \sum_j \pi_{jk} \gamma_{t-1}(j)$$

# Numerical Implementation

## Forward algorithm

$$\gamma_t(k) = p(z_t = k, x_{1:t}, y_{1:t}) = \mathcal{L}_k(y_t) \mathcal{L}_k(x_t) \sum_j \pi_{jk} \gamma_{t-1}(j)$$

For the first step

$$\gamma_1(k) = p(z_1 = k, x_1, y_1) = \mathcal{L}_k(y_1) p(x_1 | z_1 = k) p(z_1 = k)$$

Where

$p(x_1 | z_1 = k)$  is a multivariate Gaussian

$p(z_1 = k)$  is a uniform distribution



# Numerical Implementation

## Viterbi algorithm

Most of the steps are similar to the forward algorithm so we won't get in too much detail

$$\eta_t(k) = \arg \max p(z_{1:t-1}, z_t = k, x_{1:t}, y_{1:t})$$

We can retrieve the recursive relation

$$\eta_t(k) = \mathcal{L}_k(y_t) \mathcal{L}_k(x_t) \operatorname{argmax}_j [\pi_{jk} \eta_{t-1}(j)]$$

With initialization

$$\eta_1(k) = p(z_1 = k, x_1, y_1)$$

# Numerical Implementation

## Extracting information

Putting everything together we have

$$p(x_{1:T}, y_{1:T}) = \sum_{k=1}^K p(z_T = k, x_{1:T}, y_{1:T}) = \sum_{k=1}^K \gamma_T(k)$$

$$\hat{z}_{1:T} = \arg \max_{z_{1:T}} p(z_{1:T}, x_{1:T}, y_{1:T}) = \arg \max_k \eta_T(k)$$

# Numerical Implementation

## JAGS


In JAGS most of the issues that plagued the Stan implementation were, fortunately, absent and as a result the model was easier to code since:

- 1 - It is enough to specify the stochastic dependance between variables to directly retrieve posterior distributions for all parameters
- 2 - The Viterbi algorithm wasn't needed thanks to the native support for discrete parameters

# Numerical Implementation

## JAGS

On the other hand, JAGS does not allow redefining variables, resulting in the need for much more complex, multi-dimensional arrays.



Moreover, some functions are not vectorized

# Generated data

## Dummy Random Walk

The first dataset we generated to test the system was a polar random walk with two states, also known in the literature as “Markov windshield wiper”.

state 1: “move clockwise with step  $\theta$ ”

state 2: “move counter clockwise with step  $\theta$ ”


As for the transition matrix it simply was

$$\pi = \begin{pmatrix} 0.7 & 0.3 \\ 0.3 & 0.7 \end{pmatrix}$$

# Generated data

## Results

The dataset was then saved as  $(x, y)$  coordinates



# Recursive Model

## Refinement


One way to make the model more flexible is to modify the Markov process by introducing a conditional dependence on the previous state

$$z_t \sim \pi_{SB}(\nu_t) \quad \text{with} \quad \nu_t = R_{z_{t-1}} y_{t-1} + r_{z_{t-1}}$$

To have consistent dimensionality

$$R_k \in \mathbb{R}^{K-1 \times M}$$

$$r_k \in \mathbb{R}^{K-1}$$



# Recursive Model

## Stick Breaking

The stick breaking distribution is defined as

$$\pi_{SB}(\nu) = [\pi_{SB}^{(1)}(\nu), \dots, \pi_{SB}^{(K)}(\nu)]$$
$$\pi_{SB}^{(K)}(\nu) = \begin{cases} \sigma(\nu_k) \prod_{j=1}^k (1 - \sigma(\nu_j)) & \text{if } k \leq K - 1, \\ \prod_{j=1}^{K-1} (1 - \sigma(\nu_j)) & \text{if } k = K, \end{cases}$$

Where  $\sigma(x)$  is the logistic function

$$\sigma(x) = \frac{e^x}{1 + e^x}$$



# Recursive Model

## Stick Breaking

Both the Forward and Viterbi algorithms can be adapted to accommodate the Stick Breaking update rule, we will also need additional priors.

$$(R_k, r_k) \sim \text{MN}(M_r, \Sigma_r, \Omega_r)$$



# Racing circuit datasets

## NASCAR and Monza circuits

A simple code has been implemented to simulate the driving behavior around a racing track.

First a simpler NASCAR style circuit, then a heavily romanticized rendition of the Monza circuit\*


\*which notably lacks its only decent corner, the Ascari chicane. [editor note]




# Results

## NASCAR - JAGS


Final simulations have been run with 20000 burn-in steps and then generated 10 sequences of reconstructed data



Original



Reconstructed 1




Reconstructed 2


# Results

## Monza - JAGS

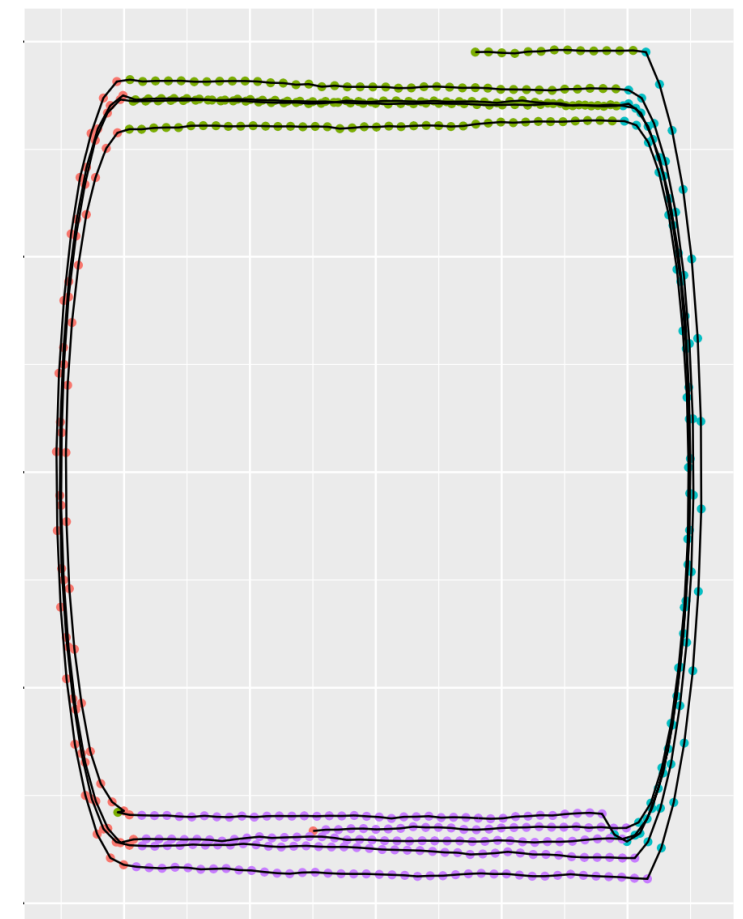
In this case the added complexity of the model tends to push the boundaries of the system, as a matter of fact stability is hard to achieve



Original



Reconstructed 1



Reconstructed 2

$$\begin{aligned}
\gamma_t(k) &= \sum_{j=1}^K p(z_t = k, z_{t-1} = j, x_{1:T}, y_{1:T}) = \\
&= p(y_t | z_t = k, x_t) p(x_t | z_t = k, x_{t-1}) \cdot \sum_j \pi_{jk} p(z_{t-1} = j, x_{1:t-1}, y_{1:t-1}) =
\end{aligned}$$